

```
In [ ]: ##### Author : Amir Shokri
##### github link : https://github.com/amirshnll/Abalone
##### dataset link : http://archive.ics.uci.edu/ml/datasets/Abalone
##### email : amirsh.nll@gmail.com
```

```
In [2]: import sklearn
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [3]: #read file
df = pd.read_csv("D:\\abalone.txt", header=None)
for char in df:
    df = df.replace('M', '1')
    df = df.replace('F', '-1')
    df = df.replace('I', '0')
df

#separate the feature columns from the target column.
features = [0,1,2,3,4,5,6,7]
X = df[features]
y = df[8]
print(X)
print(y)
```

	0	1	2	3	4	5	6	7
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700
2	-1	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550
4	0	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550
...	..	...	...	...	...	...	...	...
4172	-1	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490
4173	1	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605
4174	1	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080
4175	-1	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960
4176	1	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950

```
[4177 rows x 8 columns]
0      15
1       7
2       9
3      10
4       7
...
4172   11
4173   10
4174    9
4175   10
4176   12
Name: 8, Length: 4177, dtype: int64
```

```
In [4]: #separate the Training data and Test data
X_train, X_test, y_train, y_test = train_test_split(X,y,random_state=1, test_size=0.2)
# Feature scaling
scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [5]: # Finally for the Logistic Regression
logisticRegr = LogisticRegression(solver='newton-cg', random_state=0 ,max_iter=
2000)
logisticRegr.fit(X_train, y_train.values.ravel())
```

```
Out[5]: LogisticRegression(max_iter=2000, random_state=0, solver='newton-cg')
```

```
In [6]: #In the prediction step, the model is used to predict the response for given d
ata.
predictions = logisticRegr.predict(X_test)
print(predictions)
```

```
[ 8  8  7  8  9  6 11  9 11  8 10  5  9 11  7  9 10 11 10  9  8 10  7 10
  9 11 11  9 11 11 10  9 10  9  6  8 10 13 13 10 10 11 11 10  8  8 11  9
 10 10  9  8 10  5 11  9  9  9  8  7  8  8 11  9  8  8 10 10  9  8  9 11
  6  9  8  7 11 10  9 10  7  8 10  8  9  8 10  7 11  9 13  8  8 11  6  8
  9 12  6 10  8 10  7  8 13  8 10  9 11  9  7  9  9  7  6 10  8  7 10 11
  9  7  7 13  8 11 10 10  8 10  8  7  8  8  9  7  9 13  8  9  7  9  8 10
  8  9 10 11  9  8 10  6  9 10  8  9 10  9  9 11  9 10 13  6  8  8  9 10
  8 13  6 10 11  8  8 10 10 10  7  8 10  9  8  7  7  9 11  7 10 11  6  9
  6 11  8 11  7  4  7  9  4  9 11  8  8  8  9  8  8  7 10  8 10  8  8  8
 10  9 11  9  8  8 10  9  8  9 10  9 10  9  9  8 10 11  8  5 11 10 10  9
 13 10  9  9  7 11 10  7  8 10 10 11 11 10  5 13 10 11  7 10 10 10  7  9
  8  9  9  9  8  9  8 13 11  8 11 13  7  8 10 11  9  8 10  6  6  7  5  6
  6  9  7  8 10  9  8  9  7 10 10  9 11  8  8 11 11  9 13  8  8  8  8 11
 10  8  7 16 11 11  8  5 10  9 10 10  8 11  8  9  9  9 10  8 10  7  8  8
 10  4 11 16  6  7 11  9  9 10 18 10 10  9  9  4  9  8  5  9  8  9 13  9
  9  5  9  9  8 10  6  8 11  6  9 10 11 10  9  9  8  9  8  4  9  7  6 13
 11  7  8 10  8  9  7 10 10  8  9  7  8  8  9 16  9  7  9  9 10  7  9 13
  8 10  9 10  9  8 16  7  9  9  8  8  9  9 11  8 13  4  6 13  9 10 10  8
 10  8  8 10  9  8  8 10 11  8  8  9 10  7  9 11  8  8 11  9  4  7  8 11
  9  8 11 13 10  7  9 11  8 11 10 11 10  9 16  9  6  8 10 10  9 10 11  4
  7  7 10 10  7  8  7  9  9 11  8  9  8  7  5  7  7  9  9  8 11  6  9  8
 11 11  9  9  9 10 16  9 10  7  6  7  7 10  9  8  8 16  9 13  8 10  9  5
 11 10  9 11  9 11 10 10  9 11  7  7  9 11  9  9  6  9  8  8  9  9  9 11
  9 11 10  8  8 10  8  8  9  7  8  9  8 10  9 10 10  8 10  7 10  8 10  8
  8 11  8  9 11  9 10  4  8 12  8  9  4 13  9 18  8 10  8  6  9 10  4 10
  9  8 10  8  7 10 11 11 10 11  8 11  9  7 11  7  7  9  9  8  9 11  7  9
  9  9  8  9  8  7  6  8  9 10  9  6 11  6  8  9  9  5  9  8  8 11  7 11
  8  7 13  9 10 11  8 13  8  9 11  7 10  9  8  9  9  9  9  7  5  6  9  9
  6  7  9  9  9 10  7 10  6  9  9  7  8  7 10  7  8 10 10  9 11  9  8  9
  7  8  9  5 11  7  7 11 10 11  9  9  7  7  7  9  7 18  9 10  8 16  8  9
  8 10 10  9  9 11 11  9 10  8  9 13  8  8  9 10 10  9 10  8  9  8  9  9
 10 10  8 10  7  8  7 10 10  7  9  8  7  7 10  5  7  7  8  8 10  7  9  8
  8 10  9  7 10 11 13  7  7  8  8  8  9 11  8  9 11 11 10 10 10 13  9 13
  9 10  8 11 10  5  7 10 10 10 11 10  7  7  9 11 10  9  7  8  9  6  9  6
  8  9  7  9 13 13  7 10 11  6  9 10  8  6 11  9 11  9 11 11]
```

```
In [7]: # Last thing: evaluation of algorithm performance in classifying  
cnf_matrix = confusion_matrix(y_test, predictions)  
print(confusion_matrix(y_test, predictions))  
print(classification_report(y_test, predictions))
```

```

[[ 0  0  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  6  4  0  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  1  8  6  5  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  1  2 16 23  9  2  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  1  9 27 36  6  1  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  1 18 46 32  7  2  0  0  0  0  1  0  0  0  0  0]
 [ 0  0  0  0  1  9 38 51 26  7  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  2  9 20 49 36 24  1  1  0  0  1  0  0  0  0  0]
 [ 0  0  0  0  0  0 11 26 28 24  0  2  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  3 10 14 16 14  0  5  0  0  1  0  0  0  0  0]
 [ 0  0  0  0  0  2  1  9  9 14  1  3  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  1  4  8  5  0  5  0  0  1  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  4 10  4  0  3  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  3  2  1  0  2  0  0  1  0  1  0  0  0]
 [ 0  0  0  0  0  0  0  3  4  2  0  2  0  0  0  0  1  0  0  0]
 [ 0  0  0  0  0  0  2  0  5  3  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  1  0  3  0  0  1  0  0  2  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  1  0  3  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  1  0  0  0  0  1  0  1  0  0  0]
 [ 0  0  0  0  0  0  0  0  1  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  2  0  0  0  0  0  0  0  0]]

```

```

precision    recall  f1-score   support

```

```

 2          0.00          0.00          0.00          1
 3          0.00          0.00          0.00          2
 4          0.55          0.55          0.55         11
 5          0.53          0.40          0.46         20
 6          0.46          0.30          0.36         53
 7          0.28          0.34          0.31         80
 8          0.26          0.43          0.33        107
 9          0.25          0.39          0.30        132
10          0.23          0.25          0.24        143
11          0.24          0.26          0.25         91
12          0.00          0.00          0.00         63
13          0.10          0.08          0.09         39
14          0.00          0.00          0.00         24
15          0.00          0.00          0.00         21
16          0.12          0.10          0.11         10
17          0.00          0.00          0.00         12
18          0.00          0.00          0.00         10
19          0.00          0.00          0.00          7
20          0.00          0.00          0.00          4
21          0.00          0.00          0.00          3
22          0.00          0.00          0.00          1
23          0.00          0.00          0.00          2

```

```

accuracy                0.26        836
macro avg              0.14        0.14        0.14        836
weighted avg          0.22        0.26        0.23        836

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\\_classification.p  
y:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and bei  
ng set to 0.0 in labels with no predicted samples. Use `zero\_division` parame  
ter to control this behavior.

```

_warn_prf(average, modifier, msg_start, len(result))

```

